

▼ Simple Linear Regression

▼ Importing the libraries

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import math
```

▼ Importing the dataset

```
from google.colab import drive
drive.mount('/content/drive')
```

📁 Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.m



```
dataset = pd.read_csv("/content/drive/MyDrive/Datasets/weight-h
print(dataset)
```

	Gender	Height	Weight
0	Male	73.847017	241.893563
1	Male	68.781904	162.310473
2	Male	74.110105	212.740856
3	Male	71.730978	220.042470
4	Male	69.881796	206.349801
...
9995	Female	66.172652	136.777454
9996	Female	67.067155	170.867906
9997	Female	63.867992	128.475319
9998	Female	69.034243	163.852461
9999	Female	61.944246	113.649103

[10000 rows x 3 columns]

```
X = dataset.iloc[:,1:2]
y = dataset.iloc[:, -1]
```

```
print(X)
```

	Height
0	73.847017

```

1      68.781904
2      74.110105
3      71.730978
4      69.881796
...
9995   66.172652
9996   67.067155
9997   63.867992
9998   69.034243
9999   61.944246

```

```
[10000 rows x 1 columns]
```

```
print(y)
```

```

0      241.893563
1      162.310473
2      212.740856
3      220.042470
4      206.349801
...
9995   136.777454
9996   170.867906
9997   128.475319
9998   163.852461
9999   113.649103
Name: Weight, Length: 10000, dtype: float64

```

▼ Splitting the dataset into the Training set and Test set

```

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_

```

▼ Training the Simple Linear Regression model on the Training set

```

from sklearn.linear_model import LinearRegression

regressor = LinearRegression()
# Upto here was the builing part, now we have to train the mode

# To connect our model we use the fit method

# X_train contains the features of the dataset
# y_train contins the dependant varaibles
regressor.fit(X_train, y_train)

```

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

▼ Predicting the Test set results

```
regressor.predict(X_test)
# This returns a vector containg the predicted data


array([148.7894772 , 168.43520123, 224.31884497, ..., 159.17576427,
       155.86404539, 144.83449257])
```

```
y_pred = regressor.predict(X_test)
```

▼ Visualising the Training set results

```
plt.title("Weight vs Height(Training Set)")
plt.xlabel("Height")
plt.ylabel("Weight")
# scatter method allows us to put points/coordinates
plt.scatter(X_train, y_train, color="red") # Real vlaues
```


```
<matplotlib.collections.PathCollection at 0x7fc9687762d0>
```



```
# plot method is used to plot the curve of a function(y=b0+b1*x)
plt.title("Weight vs Height(Training Set)")
plt.xlabel("Height")
plt.ylabel("Weight")
```


```
plt.plot(X_train, regressor.predict(X_train)) # Best fit line o
```

```
[<matplotlib.lines.Line2D at 0x7fc96aa54490>]
```



```
plt.title("Weight vs Height(Training Set)")
plt.xlabel("Height")
plt.ylabel("Weight")
plt.scatter(X_train, y_train, color="red")
plt.plot(X_train, regressor.predict(X_train))
# show is used to display the graphic in the output
plt.show()
```

Red dots are real values of salary (x_train, y_train)
 # Blue line is the best fit line on training X_train and predic




▼ Visualising the Test set results

```
plt.title("Weight vs Height(Training Set)")
```

```
plt.xlabel("Height")
plt.ylabel("Weight")
plt.scatter(X_test, y_test, color="red")
```

```
# Predicted salries of test set will on the same rgression line
plt.plot(X_train, regressor.predict(X_train), color="blue")
plt.show()
```

```
# Red are new observations(test set)
# Blue is our best fit line after traning on the available trai
```



Making a single prediction (for example the salary of an employee with 12 years of experience)

```
input_value = int(input("Enter Height: "))
print(regressor.predict([[input_value]]))
```

```
Enter Height: 185
[1076.04180898]
```

Therefore, our model predicts that the salary of an employee with 12 years of experience is \$ 138967,5.

Important note: Notice that the value of the feature (12 years) was input in a double pair of square brackets. That's because the "predict" method always expects a 2D array as the format of its inputs. And putting 12 into a double pair of square brackets makes the input exactly a 2D array. Simply put:

12 → scalar

`[12] → 1D array`

`[[12]] → 2D array`

Getting the final linear regression equation with the values of the coefficients

```
print(regressor.coef_)
print(regressor.intercept_)
```

```
[ 7.70936331]
[-350.1904028560757]
```

Therefore, the equation of our simple linear regression model is:

$$\text{Weight} = 7.70936331 \times \text{Height} + -350.190402856075719.$$

Important Note: To get these coefficients we called the "coef_" and "intercept_" attributes from our regressor object. Attributes in Python are different than methods and usually return a simple value or an array of values.

```
from sklearn.metrics import mean_squared_error
```

```
mse = mean_squared_error(y_test, y_pred)
print("MSE: ", mse)
```

```
MSE: 146.53677213957428
```

```
rmse = math.sqrt(mse)
print(rmse)
```

```
12.105237384684957
```

