

Name: **Aditya Kangune**

Roll number: **33323**

Batch: **K11**

▼ Multiple Linear Regression


▼ Importing the libraries

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
```

▼ Importing the dataset

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.m



```
dataset = pd.read_csv('/content/drive/MyDrive/Datasets/temperat
print(dataset)
```

	YEAR	JAN	FEB	MAR	...	JAN-FEB	MAR-MAY	JUN-SEP	OCT-DEC
0	1901	22.40	24.14	29.07	...	23.27	31.46	31.27	27.25
1	1902	24.93	26.58	29.77	...	25.75	31.76	31.09	26.49
2	1903	23.44	25.03	27.83	...	24.24	30.71	30.92	26.26
3	1904	22.50	24.73	28.21	...	23.62	30.95	30.66	26.40
4	1905	22.00	22.83	26.68	...	22.25	30.00	31.33	26.57
..
112	2013	24.56	26.59	30.62	...	25.58	32.58	31.33	27.83
113	2014	23.83	25.97	28.95	...	24.90	31.82	32.00	27.81
114	2015	24.58	26.89	29.07	...	25.74	31.68	31.87	28.27
115	2016	26.94	29.72	32.62	...	28.33	34.57	32.28	30.03
116	2017	26.45	29.46	31.60	...	27.95	34.13	32.41	29.69

[117 rows x 18 columns]

```
dataset.head()
```

	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV	DEC
0	1901	22.40	24.14	29.07	31.91	33.41	33.18	31.21	30.39	30.47	29.97	27.31	24.40
1	1902	24.93	26.58	29.77	31.78	33.73	32.91	30.92	30.73	29.80	29.12	26.31	24.00
2	1903	23.44	25.03	27.83	31.39	32.91	33.00	31.34	29.98	29.85	29.04	26.08	23.60
3	1904	22.50	24.73	28.21	32.02	32.64	32.07	30.36	30.09	30.04	29.20	26.36	23.60
4	1905	22.00	22.83	26.68	30.01	33.32	33.25	31.44	30.68	30.12	30.67	27.52	23.80

```
# X = dataset.iloc[:, :-1].values
# y = dataset.iloc[:, -1].values
```

```
dataset.dtypes
```

```
YEAR          int64
JAN           float64
FEB           float64
MAR           float64
APR           float64
MAY           float64
JUN           float64
JUL           float64
AUG           float64
SEP           float64
OCT           float64
NOV           float64
DEC           float64
ANNUAL        float64
JAN-FEB       float64
MAR-MAY       float64
JUN-SEP       float64
OCT-DEC       float64
dtype: object
```

```
dataset.columns
```

```
Index(['YEAR', 'JAN', 'FEB', 'MAR', 'APR', 'MAY', 'JUN', 'JUL', 'AUG', 'SEP',
      'OCT', 'NOV', 'DEC', 'ANNUAL', 'JAN-FEB', 'MAR-MAY', 'JUN-SEP',
      'OCT-DEC'],
      dtype='object')
```

```
dataset.describe()
```


	YEAR	JAN	FEB	MAR	APR	MAY	J
count	117.000000	117.000000	117.000000	117.000000	117.000000	117.000000	117.0000
mean	1959.000000	23.687436	25.597863	29.085983	31.975812	33.565299	32.7742
std	33.919021	0.834588	1.150757	1.068451	0.889478	0.724905	0.6331
min	1901.000000	22.000000	22.830000	26.680000	30.010000	31.930000	31.1000
25%	1930.000000	23.100000	24.780000	28.370000	31.460000	33.110000	32.3400
50%	1959.000000	23.680000	25.480000	29.040000	31.950000	33.510000	32.7300
75%	1988.000000	24.180000	26.310000	29.610000	32.420000	34.030000	33.1800

```
dataset.isnull().sum()
```

```
YEAR      0
JAN       0
FEB       0
MAR       0
APR       0
MAY       0
JUN       0
JUL       0
AUG       0
SEP       0
OCT       0
NOV       0
DEC       0
ANNUAL    0
JAN-FEB   0
MAR-MAY   0
JUN-SEP   0
OCT-DEC   0
dtype: int64
```

```
top_10_data = dataset.nlargest(10, "ANNUAL")
plt.figure(figsize=(14,12))
plt.title("Top 10 temperature records")
sns.barplot(x=top_10_data.YEAR, y=top_10_data.ANNUAL)
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f2bf41ec410>



▼ Taking in the required month for prediction

```

| ■ ■ ■ ■ ■ ■ ■ ■ ■ ■
month_input = input("Enter month for prediction (Eg:'JAN'):")

```

```

Enter month for prediction (Eg:'JAN'):APR

```

```

| ■ ■ ■ ■ ■ ■ ■ ■ ■ ■
X = dataset[month_input].values
y = dataset["ANNUAL"].values

```

```

# X = np.array(X, ndmin=2)
# y = np.array(y, ndmin=2)

```

```

print(X)

```

```

[31.91 31.78 31.39 32.02 30.01 31.93 31.79 32.42 30.79 31.42 31.27 31.29
 32.02 30.96 31.36 31.99 30.61 30.68 31.55 30.89 32.47 31.85 32.37 32.07
 32.53 30.42 31.5 31.7 31.59 30.98 32.65 32.08 30.53 31.95 30.24 31.52
 30.96 32.33 31.03 31.38 32.8 31.98 30.75 30.93 30.89 32.19 32.28 31.93
 31.85 31.16 30.3 32.22 32.1 32.47 31.12 32.04 30.91 32.51 32.33 31.9
 31.8 31.76 30.99 32.12 30.6 31.95 31.7 31.59 32.4 32.98 32.19 31.46
 33.3 32.68 32.36 31.71 31.7 32.23 33.11 33.36 32.26 31.65 31.1 32.19
 32.72 32.11 32.18 32.15 31.76 31.59 31.51 31.89 32.33 31.57 32.83 32.4]

```

```
31.    32.6   33.77  33.17  32.54  33.51  32.91  32.97  32.37  32.59  33.57  32.13
22  00  21  07  21  7   22  16  22  66  22  71  21  07  25  20  21  051
```

```
print(y)
```

```
[28.96 29.22 28.47 28.49 28.3   28.73 28.65 28.83 28.38 28.53 28.62 28.95
 28.67 28.66 28.94 28.82 28.11 28.66 28.66 28.76 28.86 28.8   28.74 28.8
 28.67 28.7   28.59 28.98 28.76 28.65 29.15 29.09 28.49 29.03 28.76 28.71
 28.7   28.7   28.85 28.88 29.46 28.98 28.8   28.89 28.97 29.37 28.84 28.73
 28.89 28.47 29.09 29.16 29.43 28.92 28.76 28.63 28.64 29.34 29.02 29.31
 28.72 28.89 29.04 29.09 29.16 29.41 29.14 29.07 29.61 29.47 29.15 29.31
 29.44 29.26 28.89 29.27 29.41 29.23 29.63 29.58 29.32 29.12 29.11 29.28
 29.61 29.33 29.72 29.55 29.18 29.14 29.32 29.23 29.55 29.46 30.18 29.58
 29.05 29.7   29.81 29.75 29.99 30.23 29.75 29.79 29.6   30.06 29.84 29.64
 30.3   30.13 29.82 29.81 29.81 29.72 29.9   31.63 31.42]
```

▼ Splitting the dataset into the Training set and Test set

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_
```

▼ Training the Linear Regression model on the Training set

```
X_train= X_train.reshape(-1, 1)
y_train= y_train.reshape(-1, 1)
X_test = X_test.reshape(-1, 1)
```

```
from sklearn.linear_model import LinearRegression
regressor = LinearRegression()
regressor.fit(X_train, y_train)
```

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

▼ Predicting the Test set results

```
y_pred = regressor.predict(X_test)
# np.set_printoptions(precision=2)
# print(np.concatenate((y_pred.reshape(len(y_pred),1), y_test.r

# print(y_test.reshape()))

print(y_test)
```

```
[28.62 29.31 29.58 29.23 28.83 29.72 28.59 29.81 28.74 30.18 30.23 28.47
 29.09 28.67 31.42 29.04 29.46 28.89 28.89 29.26 28.11 30.3 28.66 28.89]
```

```
print(y_pred)
```

```
[[28.85675571]
 [29.13776779]
 [29.36079324]
 [29.13330728]
 [29.36971426]
 [29.26266204]
 [28.95934742]
 [29.47676648]
 [29.34741171]
 [29.55259513]
 [29.85590975]
 [28.91028182]
 [28.42408633]
 [29.41877986]
 [30.49822306]
 [28.73186146]
 [28.99057099]
 [29.3429512 ]
 [29.07532066]
 [29.48568749]
 [28.56236211]
 [29.66856837]
 [28.71847993]
 [28.70509841]]
```

```
y_test = np.array(y_test)
y_pred = np.array(y_pred)
```

▼ Mean Squared Error

```
from sklearn.metrics import mean_squared_error
mse = mean_squared_error(y_test, y_pred)
print(mse)
```

```
0.21199560831993922
```

▼ Root Mean Squared Error

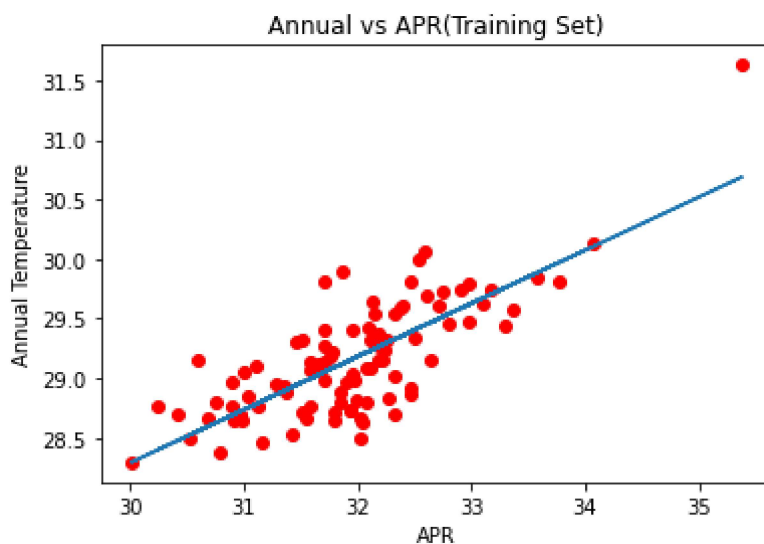
```
import math
rmse = math.sqrt(mse)
print(rmse)
```

0.4604298082443612

▼ Visualising the Training set results

```
plt.title("Annual vs {month}(Training Set)".format(month=month_
plt.ylabel("Annual Temperature")
plt.xlabel("{month}".format(month=month_input))
plt.scatter(X_train, y_train, color="red")
plt.plot(X_train, regressor.predict(X_train))
# show is used to display the graphic in the output
plt.show()
```

Red dots are real values of salary (x_train, y_train)
 # Blue line is the best fit line on training X_train and predic




▼ Visualising the Test set results

```
plt.title("Annual Temperature vs {month}(Testing Set)".format(m
plt.ylabel("Annual Temperature")
plt.xlabel("{month}".format(month=month_input))
plt.scatter(X_test, y_test, color="red")
```

Predicted salaries of test set will on the same rgression line

```
plt.plot(X_train, regressor.predict(X_train), color="blue")
plt.show()
```

```
# Red are new observations(test set)
# Blue is our best fit line after traning on the available trai
```



```
print(regressor.coef_)
print(regressor.intercept_)
```

```
[[0.44605091]]
[14.90874382]
```

Therefore, the equation of our simple linear regression model is:

$$\text{Annual_temperature} = 0.44605091 \times \text{month} + 14.90874382.$$

Important Note: To get these coefficients we called the "coef_" and "intercept_" attributes from our regressor object. Attributes in Python are different than methods and usually return a simple value or an array of values.

