

Operating Systems Lab

Assignment 9

Name: Aditya Kangune

Roll number: 33323

Batch: K11

Aim:

Implement a new system call, add this new system call in the Linux kernel (any kernel source, any architecture and any Linux kernel distribution) and demonstrate the use of same.

Objective:

Add a new system call, `swipe()`, to the Linux kernel that transfers the remaining time slice of each process in a specified set to a target process. You will also demonstrate various uses of the system call (both advantageous and detrimental)

Theory:

Adding a simple system call:

1. Download the kernel source:

In your terminal type the following command:

```
wget https://www.kernel.org/pub/linux/kernel/v4.x/linux-4.17.4.tar.xz
```

Else go to [kernel.org](https://www.kernel.org) and download the latest version.

2. Extract the kernel source code:

```
sudo tar -xvf linux-4.17.4.tar.xz -C/usr/src/
```

tar — Tar stores and extracts files from a tape or disk archive.

-x — extract files from an archive

-v — requested using the -verbose option, when extracting archives

-f — file archive; use archive file or device archive

-C — extract to the directory specified after it.(in this case /usr/src/) Now, we'll change the directory to where the files are extracted:

3. Define a new system call sys_hello():

Create a directory named hello/ and change the directory to hello/:

```
mkdir
```

```
hello cd
```

```
Hello
```

Create a file **hello.c** using text editor:

```
gedit hello.c
```

Write the following code in the editor:

```
#include <linux/kernel.h>
asmlinkage long sys_hello(void)
{
    printk(KERN_INFO "Hello world\n");
    return 0;
}
```

printk prints to the kernel's log file.

Create a “Makefile” in the hello directory:

gedit Makefile

and add the following line to it:

obj-y := hello.o

This is to ensure that the hello.c file is compiled and included in the kernel source code.

Note: There is no space in between “obj-y”.

4. Adding hello/ to the kernel's Makefile:

Go back to the parent dir i.e. `cd ../` and open “Makefile”

gedit Makefile

search for core-y in the document, you'll find this line as the second instance of your search:

core-y += kernel/ mm/ fs/ ipc/ security/ crypto/ block/

Add 'hello/' to the end of this line:

core-y += kernel/ mm/ fs/ ipc/ security/ crypto/ block/ hello/

Note: There is a space between “block/” and “hello/”. (Doing such a mistake may cause errors in further steps)

This is to tell the compiler that the source files of our new system call (`sys_hello()`) are in present in the hello directory.

5. Add the new system call to the system call table:

If you are on a 32-bit system you'll need to change 'syscall_32.tbl'.

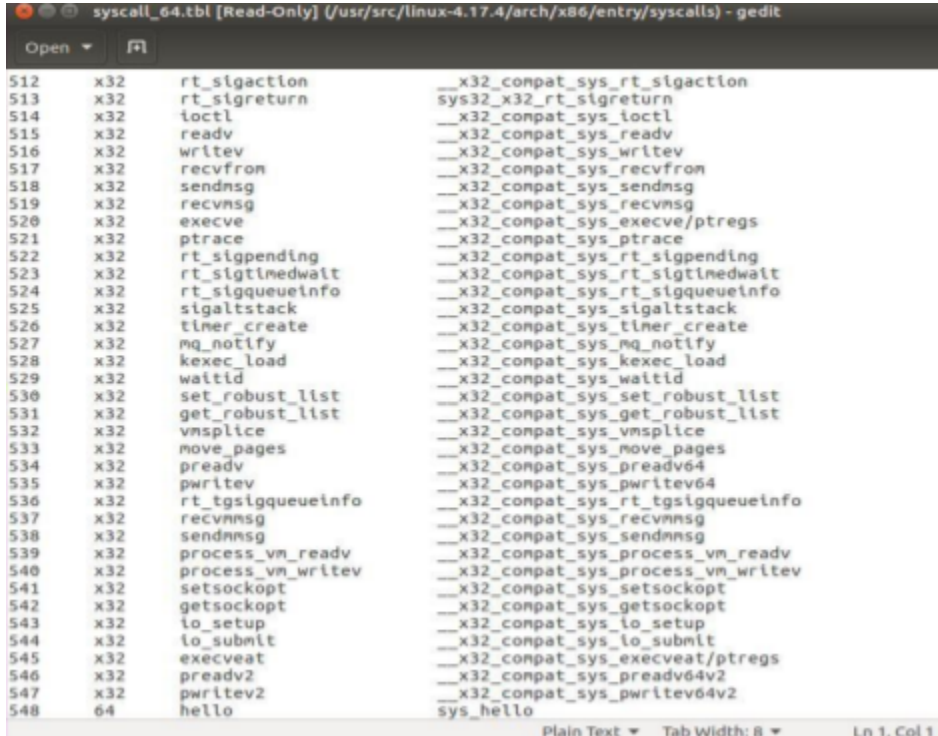
For 64-bit, change 'syscall_64.tbl'.

Run the following commands in your terminal from linux-4.17.4/ directory:

cd arch/x86/entry/syscalls

/ gedit syscall_64.tbl

You'll get a file like the following in your editor:



```
512 x32 rt_sigaction __x32_compat_sys_rt_sigaction
513 x32 rt_sigreturn sys32_x32_rt_sigreturn
514 x32 ioctl __x32_compat_sys_ioctl
515 x32 readv __x32_compat_sys_readv
516 x32 writev __x32_compat_sys_writev
517 x32 recvfrom __x32_compat_sys_recvfrom
518 x32 sendmsg __x32_compat_sys_sendmsg
519 x32 recvmmsg __x32_compat_sys_recvmmsg
520 x32 execve __x32_compat_sys_execve/ptregs
521 x32 ptrace __x32_compat_sys_ptrace
522 x32 rt_sigpending __x32_compat_sys_rt_sigpending
523 x32 rt_sigtimedwait __x32_compat_sys_rt_sigtimedwait
524 x32 rt_sigqueueinfo __x32_compat_sys_rt_sigqueueinfo
525 x32 sigaltstack __x32_compat_sys_sigaltstack
526 x32 timer_create __x32_compat_sys_timer_create
527 x32 mq_notify __x32_compat_sys_mq_notify
528 x32 kexec_load __x32_compat_sys_kexec_load
529 x32 waitid __x32_compat_sys_waitid
530 x32 set_robust_list __x32_compat_sys_set_robust_list
531 x32 get_robust_list __x32_compat_sys_get_robust_list
532 x32 vmsplice __x32_compat_sys_vmsplice
533 x32 move_pages __x32_compat_sys_move_pages
534 x32 preadv __x32_compat_sys_preadv64
535 x32 pwritev __x32_compat_sys_pwritev64
536 x32 rt_tgsigqueueinfo __x32_compat_sys_rt_tgsigqueueinfo
537 x32 recvmmsg __x32_compat_sys_recvmmsg
538 x32 sendmmsg __x32_compat_sys_sendmmsg
539 x32 process_vm_readv __x32_compat_sys_process_vm_readv
540 x32 process_vm_writev __x32_compat_sys_process_vm_writev
541 x32 setsockopt __x32_compat_sys_setsockopt
542 x32 getsockopt __x32_compat_sys_getsockopt
543 x32 io_setup __x32_compat_sys_io_setup
544 x32 io_submit __x32_compat_sys_io_submit
545 x32 execveat __x32_compat_sys_execveat/ptregs
546 x32 preadv2 __x32_compat_sys_preadv64v2
547 x32 pwritev2 __x32_compat_sys_pwritev64v2
548 64 hello sys_hello
```

Go to the last of the document and add a new line like so:

548 64 hello sys_hello

Note: Here 548 is written because in the previous line the number entry was 547. Remember this number it will be used in the later steps.

Also, note that I've written 64 in my system because it is 64 bit. You may have to write i586 or x32. For knowing what is to be written check in this file itself in many of the lines you may find entries like so:

308	common	setns	__x64_sys_setns
309	common	getcpu	__x64_sys_getcpu
310	64	process_vm_readv	__x64_sys_process_vm_readv
311	64	process_vm_writev	__x64_sys_process_vm_writev
312	common	kcmp	__x64_sys_kcmp
313	common	finit_module	__x64_sys_finit_module
314	common	sched_setattr	__x64_sys_sched_setattr
315	common	sched_getattr	__x64_sys_sched_getattr
316	common	renameat2	__x64_sys_renameat2
317	common	seccomp	__x64_sys_seccomp
318	common	getrandom	__x64_sys_getrandom
319	common	memfd_create	__x64_sys_memfd_create
320	common	kexec_file_load	__x64_sys_kexec_file_load
321	common	bpf	__x64_sys_bpf
322	64	execveat	__x64_sys_execveat/ptregs

64 written at 310, 311 and 322 line numbers.

This will tell you whether to write i586 or something else.

Save and exit.

6. Add new system call to the system call header file:

Go to the linux-4.17.4/ directory and type the following commands:

cd include/linux/

gedit syscalls.h

Add the following line to the end of the document before the #endif statement:

asmlinkage long sys_hello(void);

Save and exit. This defines the prototype of the function of our system call. "asmlinkage" is a key word used to indicate that all parameters of the function would be available on the stack.

7. Compile the kernel:

Before starting to compile you need to install a few packages. Type the following commands in your terminal:

`sudo apt-get install gcc`

```
sudo apt-get install
```

```
libncurses5-dev sudo apt-get
```

```
ev sudo apt-get
```

```
install bison
```

```
sudo apt-get install flex
```

```
sudo apt-get install
```

```
libssl-dev sudo apt-get
```

```
install libelf-dev sudo
```

```
apt-get update
```

```
sudo apt-get upgrade
```

to configure your kernel use the following command in your **linux-4.17.4/** directory:

sudo make menuconfig

Once the above command is used to configure the Linux kernel, you will get a pop up window with the list of menus and you can select the items for the new configuration. If your unfamiliar with the configuration just check for the file systems menu and check whether “ext4” is chosen or not, if not select it and save the configuration.

Now to compile the kernel you can use the make command:

sudo make

Pro Tip:

The make command can take a lot of time in compiling, to speed up the process you can take advantage of the multiple cores that our systems have these days.

Simply type,

sudo make -jn

where n is the number of cores that you have in your linux system.

For example if you have a Quad core(4) processor, you can write:

sudo make -j4

this will speed up my make process 4x times.

This might take an hours or more depending on your system.

8. Install / update Kernel:

Run the following command in your terminal:

```
sudo make modules_install install
```

It will create some files under /boot/ directory and it will automatically make a entry in your grub.cfg.

To check whether it made correct entry, check the files under /boot/ directory .

If you have followed the steps without any error you will find the following files in it in addition to others.

1. System.map-4.17.4

2. vmlinuz-4.17.4

3. initrd.img-4.17.4

4. config-4.17.4

Now to update the kernel in your system reboot the system.

You can use the following command:

shutdown -r now

After rebooting you can verify the kernel version using the following command:

uname -r

It will display the kernel version like so:

4.17.4

9. Test system call:

Go to your home(~) directory using the following commands and create a userspace.c file.

```
cd ~
```

```
gedit userspace.c
```

```
#include <stdio.h>
#include <linux/kernel.h>
#include <sys/syscall.h>
#include <unistd.h>
int main()
{
    long int r = syscall(358);
    printf("System call sys_hello returned %ld\n", r);
    return 0;
}
```

Note: Remember to keep in mind the number of system call that is added in syscalls_64.tbl? In my case the number was 548. Write that same number in your userspace.c file as an argument in syscall() function.

Now, compile and run the program:

```
gcc userspace.c
```

```
./a.out
```

If all the steps are done correctly you'll get an output like below:

System call sys_hello returned 0

Now, to check the message of your kernel run the following command:

dmesg

This will display Hello world at the end of the kernel's message.

Write the following code in this file:

Conclusion:

A new system call, add this new system call in the Linux kernel (any kernel source, any architecture, and any Linux kernel distribution) was studied.